

ARCQUS.AI

The LLM Wiki

Educational Primer

Explained from scratch: what it is, what it looks like, and how to put it to work.

VERSION 1.2 | OCTOBER 2026

Over the past two years, I have had tremendous success building incredibly powerful AI agents, but I've always been frustrated that an agent is only as good as the knowledge base it pulls from. Information gets stale. Data doesn't always carry forward, and a comprehensive and dynamic repository of the canonical information just isn't easy to find.

This wiki process has changed that. Now I have a dynamic, human-ratified master library of the most current thinking at my fingertips for myself and my agents. It has changed how I use the tools and how I create value from them.

Dusty Holcomb, Arcqus.AI

You do not need to know anything going in, and there is no original post you needed to read first. This guide starts at zero and builds the whole picture from the ground up: what this is, what it looks like on a real screen, what you actually do with it, and why, once you see it, you will not want to work the old way again.

One quick orientation, so an unfamiliar phrase is not hanging over you the whole way. An LLM is the kind of AI behind tools like ChatGPT and Claude. This guide is about a simple way to use one so that what you learn accumulates in a single place and gets richer over time, instead of being worked out from scratch every time you ask a question. That is the entire idea. Everything that follows just makes it concrete.

1. The Problem You Do Not Know You Have

The most valuable thing most of us do in knowledge work is synthesize. You take messy inputs, a meeting, a few articles, a long document, a hallway conversation, a half-formed hunch, and you compress them into something clear. A pattern. A connection. A position. That compression is the valuable part. It is the thing people actually pay for.

And here is what nobody points out. That synthesis is the one thing you never keep.

Look at what survives a normal week. The raw stuff survives. The recording sits in a folder, the articles are bookmarked, the notes are in a doc. But the gold, the connection you drew, the pattern you spotted across three projects, the insight that made you sit up mid-conversation, happened in your head or in a chat window, and then it was gone. You lit a fire, it gave off light, and you let it burn out. Next week you light a fresh one from scratch.

That is the problem. Not that you do not read enough or think enough. You read and think constantly. The trouble is that your thinking does not accumulate. Each insight is an event, not a deposit. Nothing compounds.

This guide is about a way to make the fire keep burning, and to make every new log catch from the last one.

2. The Idea, in One Analogy

Forget technology for a second. Here is the whole concept as a story about an intern.

The old way

You hire a brilliant intern who has no memory. Every time you have a question, they sprint to your filing cabinet, speed-read every relevant document from scratch, hand you a sharp answer, and then forget everything the instant they are done. Ask the same question next week and they re-read the entire cabinet again. They are fast and smart, but they never build a picture of your world. They never get better at your material. They start from zero every single time.

That is how almost everyone uses AI with their documents today. You upload files, the AI reads the relevant bits, answers, and forgets. You may have heard the term RAG for this. It just means the AI fetches chunks of your files at question-time. ChatGPT file uploads and NotebookLM work this way. It works. But nothing is ever built.

The new way

Same brilliant intern, but now they keep a notebook. Every source you hand them, they read once and write into the notebook: a clean summary, and, more importantly, an update to every related page, fresh cross-references, and a flag anywhere the new source contradicts something you told them before. Ask a question and they check the notebook, which already holds the synthesized answer, instead of re-reading the cabinet. They point you to exactly which sources back it up.

And here is the kicker. When you have a good insight, they write that into the notebook too. Your conclusions get saved right alongside what you read. The notebook gets richer every week, from your sources and from your own thinking. After a year, the notebook is worth more than the filing cabinet ever was, because the cabinet is raw material and the notebook is understanding.

That notebook is the wiki. That is the entire idea. Everything else is detail about how to make the intern reliable.

> *"A search system makes your documents findable. A wiki makes your thinking cumulative."*

3. The Three Pieces

There are only three moving parts. Map them straight onto the intern story.

- **Raw sources:** your filing cabinet. Articles, transcripts, PDFs, notes. The intern reads these but never changes them. They are your source of truth, and they stay exactly as they are.
- **The wiki:** the notebook. A set of pages the intern writes and maintains, one page per important topic, person, idea, or framework, all linked together. You read this; you basically never write it by hand. The intern owns it.
- **The maintainer:** the intern, which is the AI. Its whole job is the tedious part you would never keep up with yourself, reading, summarizing, cross-referencing, updating fifteen pages when one new source changes the picture, noticing contradictions. It does not get bored and it does not forget a link. That is why this works now and did not before, because the bookkeeping that always killed knowledge systems is exactly what the AI is good at.

Hold this distinction, because it is the one most people blur. The AI is not the knowledge. The wiki is the knowledge. The AI is the worker that keeps it current. If the AI vanished tomorrow, your wiki would still be sitting there, fully readable, because it is just a stack of written pages.

4. What This Looks Like on Your Screen

This is the part that stays fuzzy until you have seen it, so let us make it concrete. Picture two windows, side by side.

Left window: your AI agent (you talk to the worker)	Right window: Obsidian (you read the result)
It looks like a chat, but this AI has hands. You type an instruction, "read this new article and update the wiki," and it opens files on your hard drive, reads them, edits them, and saves them. You watch it work, opening pages, rewriting them, logging what it did.	A clean reader for your folder of text files. It shows your pages, makes the links between them clickable so you can follow a thread from one idea to the next, and offers a graph view, every page a dot and every link a line, so you can see the shape of your knowledge at a glance.

Claude Code and Codex both work. So does any coding agent that can read and edit files in a folder.

Here is the leap most newcomers have not seen: this AI has hands. The chatbot in your browser can only talk. It cannot touch your computer. Your AI agent is the same kind of intelligence, but it

can create and change files in a folder you point it at. That is the difference between an advisor who gives you advice and a worker who does the task. That difference is everything.

Obsidian is just a pleasant reader, so do not overthink it. You read there. You do not edit there. You watch the worker's changes show up.

The trick is that both windows point at the same folder. Your agent writes to it; Obsidian shows you what changed. You give instructions on the left, you read the living result on the right.

What is actually in the folder

Plain text files. That is the whole technology. No database, no special app, no cloud account. The folder holds your sources, your wiki pages written in markdown (text with light formatting like headings and bullets), and a few small housekeeping files the worker maintains for you.

- A raw sub-folder holding your sources.
- A wiki sub-folder holding the pages, one text file per topic, person, idea, or framework.
- AGENTS.md, the rulebook that tells the worker how you like things done.
- index.md, the table of contents.
- log.md, the diary of what has been done, and when.
- A _state sub-folder holding the source ledger (every source the worker has read) and the review queue (decisions waiting for you).

Because it is nothing but text files in a folder, it is yours permanently. You could open it in Notepad. You could copy it to a thumb drive. No company can lock you out of it or shut it down. The whole folder also sits in something called git, which quietly saves a version every time anything changes, so you have a complete history of how your thinking evolved, for free.

That is the entire physical reality. Two windows looking at one folder of text files. Once you have watched the worker update a page in the left window and seen it appear in the right, the fuzziness is gone for good.

5. What You Actually Do

There are four moves: ingest, query, file back, and lint. Learn them at this pace.

Ingest: fold in a new source

You drop a source into the raw folder, then tell the worker, roughly, to read this new file, give you the key takeaways first, and tell you which pages it will create or update before it writes. You read the takeaways, steer it a little, and say go. It edits the relevant pages, updates the table of contents, and writes a line in the diary. One source usually touches ten or fifteen pages,

because that is what interlinking means. A new fact might update the page for the thing it concerns, two related concept pages, and the overview all at once. If a new source only ever touches one page, you have built a folder, not a wiki.

Query: ask it a question

You ask the wiki something. The worker reads the table of contents first to find the right pages, reads them, and answers, with citations back to the pages and sources, so you can trust it. This is the step that replaces opening a blank chat to think something through. Same instinct, different door. You reach into the thing that already holds your accumulated understanding.

Lint: check yourself for rot

Every so often you ask the worker to health-check the whole wiki. Are any two pages contradicting each other, has a newer source made an old claim wrong, are there orphan pages nothing links to, is there an idea mentioned everywhere that deserves its own page. Wikis do not decay from neglect of reading. They decay from contradiction and staleness piling up invisibly. Lint is the immune system. Do it on a schedule.

File back: never let a good answer die in a chat

This is the one almost everyone skips, and it is the one that makes the whole thing compound. When you ask a question and get a strong answer, a sharp comparison, a connection you had not seen, a position you just sharpened, save it back into the wiki as a new page. Do not let it die in the chat. This is how your own thinking becomes part of the permanent record, not just the sources you fed in. The rare, valuable part of the notebook is not what you read. It is what you concluded. Filing answers back is how conclusions get kept.

6. How It Plugs Into Your Real Work

Here is the worry that is natural and also misplaced. Am I going to have to carve out a new block of time to feed and tend this thing? No. The right way to wire it in is to attach it to work you already do, never to invent a wiki hour that will not survive the month.

Every workstream you have already has two natural connection points. A drop point, the artifact it already produces, which is where feeding attaches. And a pull point, the moment you would otherwise start from a blank page, which is where drawing from it attaches.

You do not change your work. You add a five-minute step at each point. A few everyday examples:

Something you already do	Drop point (feed it here)	Pull point (ask it here)
Finish a meeting or call	The notes or transcript, ingest it	Prepping the next one or a recurring review, query that topic
Read an article or report	Clip it, ingest it	Writing something and needing the clean point, query
Research a topic over weeks	Each source, ingest it	Needing the current state of your own thinking, query
Look into a competitor or market	A note on what you learned, ingest it	Needing a fast, grounded take, query

A Tuesday, made concrete

You finish a meeting at 10am. Before the next thing, you spend five minutes. Drop the notes in raw, tell the worker to ingest them, skim the takeaways, watch the relevant page update. Done. At 2pm you sit down to write something and you would normally open a blank chat. Instead you open your agent and ask the wiki for the cleanest version of a point you keep coming back to. It answers, with links to the sources where it first came up. You write from that. The writing sharpens a phrase, so you tell the worker to file the sharpened version back. During the weekly review you already do, you run one lint pass and ask one big question. That is the whole practice. It disappears into the seams of a normal week.

Manual first, automatic later

At the start, you are the wire. You move the notes, you run the ingest. Once that rhythm is second nature, you can automate the boring step, so the worker pulls the source itself when you ask it to. But do not automate something that is not yet a habit. Walk first.

7. Why This Changes Everything

If you only take the mechanics, you have missed the point. Here is what a sharp person realizes once the picture lands.

It compounds, and almost nothing else you use does

Most tools give you the same value every time. A great template helps your fiftieth use about as much as your first. That is linear. It is a wage. You show up, you get value, you show up again. The wiki is different. Its fiftieth answer is worth far more than its first, because it stands on fifty sources and forty-nine saved insights. That is not a wage. It is an asset, and it appreciates while you sleep.

Picture the curve honestly. Month one feels underwhelming, because you are asking questions of a nearly empty notebook. Month six, it is answering things you could not have asked any other tool. Year two, it holds a body of synthesized thinking that literally no one else on earth has, because it is the compounded record of how you specifically see the world. The discipline that matters most is sticking with it through the unrewarding early weeks, when the compounding has not kicked in yet and quitting feels reasonable.

It makes a new kind of question answerable

Today you can answer retrieval questions, what did I say about that. That is a search. With a kept memory, you can answer synthesis questions, what is the pattern across everything I have done this year that I cannot see from inside any single piece of it. No blank chat can answer that. No filing cabinet can. It requires a memory larger than a single conversation, which is exactly what the wiki is.

Here is why that matters. Those synthesis questions are where your best original thinking comes from. The recurring pattern across your projects that is worth turning into a framework. The spine of a long piece of writing, surfaced from everything you have already produced. The point of view you can defend because you have tracked a field over time. Those are not things you retrieve. They are things only a kept, compounding memory can produce.

Your thinking becomes the asset

This is the heart of it. Plenty of systems can store what you consume. This one stores what you conclude. The notebook holds your synthesis, kept current, traceable to its sources, getting deeper every week. You have been generating that synthesis your whole career and letting almost all of it evaporate. This is the first time it can be kept.

It becomes a thinking partner that knows what you know

Not a clever stranger you re-brief every time, but a partner with your accumulated context already loaded, one that can say you argued the opposite of this eight months ago, here is where. That is a different relationship with a tool than you have ever had.

8. What It Would Have Meant a Year Ago

Run the counterfactual, because it makes the cost of not having this concrete.

A year ago, every insight from a meeting lived for one conversation and then sank into notes you never reopened. Every article you read informed a single task and then faded. You re-researched the same topics every few months because last time's understanding was not anywhere durable. You re-explained your own thinking to every fresh AI chat, from scratch, every time. Your best thinking was a series of bright, brief flashes, and the connections between the flashes, which is where the real value lives, were never captured at all.

With this, a year ago, you would be somewhere very different now. You would have a living map of your field or topic, with a point of view that got sharper every week instead of resetting to zero. You would have one trustworthy articulation of every idea you rely on, which means any long piece of writing you owe would be half-drafted as a byproduct of work you were doing anyway. Every recurring project or relationship would have a page that made prep a five-minute read instead of an afternoon of digging. And the patterns running across your work, the ones worth naming, would have surfaced on their own, because something was finally watching the whole picture at once.

Most of all, you would think differently about your own knowledge. You would stop treating it as something you re-summon under pressure and start treating it as something you build and own, an estate that grows, not a fire you relight each morning. That shift in stance is the real prize. The wiki is just what makes it possible.

9. The Mindset Shifts to Internalize

Five reframes. If these become reflexes, you have got it.

- Chat is throwaway; the files are the asset. Stop trying to keep conversations. Keep pages.
- You curate and ask; the AI maintains. Your job is sourcing, direction, and good questions. The grunt work, summarizing, linking, filing, is the AI's job. Do not do its job; do not let it do yours.
- The value is in reaching for it, not in building it. A beautiful wiki you never open is worth less than a thin one woven into your daily work.
- Do not trust what you cannot trace. Every claim should link to its source. A synthesis you cannot trace is one you cannot rely on when it counts.
- Do not add time; attach to what you already do. The practice lives in the seams of an existing week, or it does not live at all.

10. Your First Afternoon

Do not try to build the cathedral. Prove the loop once, small, and let your own eyes do the convincing.

- Pick one narrow topic you are already curious about, not your whole world, one slice.
- Make a private copy of the starter kit. Point your agent at it and Obsidian at it. The free starter kit on the last page sets this up for you.
- Drop three or four sources into the raw folder.
- Ingest them one at a time. Read the takeaways. Watch the pages appear in Obsidian.
- Ask the wiki one real question. Get a cited answer.
- File that answer back as a new page.
- Open the graph view and look at the little constellation you just made.

That is the complete loop, end to end, in an afternoon. The moment you watch the worker update eleven pages from one source, and then watch your own filed answer become a permanent part of the picture, the abstraction collapses into something obvious. From there, the only question left is which part of your world deserves a wiki first, and that is a good problem to have.

Glossary: The Words You Will Hear

- **LLM:** short for large language model, the kind of AI behind tools like ChatGPT and Claude. In this setup it does the reading, summarizing, cross-referencing, and filing.
- **RAG:** the common approach where the AI reads chunks of your files at question-time and then forgets. The thing this pattern improves on.
- **Markdown:** plain text with light formatting like headings and bullets. What the pages are written in, and readable forever.
- **AI coding agent:** an AI that can read and edit files in a folder on your computer, such as Claude Code or Codex. It is the worker with hands. You operate it like a chat, but it does real work on real files.
- **AGENTS.md:** the rulebook file many AI agents read automatically when they open a folder. In this kit it holds the rules that make the agent a careful wiki maintainer.
- **Obsidian:** a free app that reads a folder of markdown files nicely, with clickable links between pages plus a graph view. Your reading room.

- **Git:** a system that automatically saves a version of your folder every time it changes, giving you full history and an undo button for free.
- **Frontmatter:** a few label lines at the top of a page, like date, type, and status, that let tools sort and filter. Optional.
- **MCP:** a way to connect the AI directly to a service, like your meeting recorder, so it can pull a source itself. An automation you add later.
- **Lint:** the health-check pass that finds contradictions, stale claims, and orphan pages. The word is borrowed from programming.

The One-Line Summary

Old way: a brilliant intern who reads your filing cabinet from scratch every time and forgets. New way: the same intern keeps a notebook that gets richer with every source and every insight of yours, until the notebook, not the cabinet, is the valuable thing. The notebook is yours, it is plain text, it lives in a folder, and it compounds. Your only jobs are to feed it, ask it good questions, and never let a good answer die in a chat.

Start Here

Everything in this primer is built and waiting for you in a free starter kit.

github.com/wareagledusty/wiki-federation-starter-kit

1. Click **Use this template** and make a private copy, then put it on your computer. The kit's README shows how.

2. Open it in your AI agent. Claude Code and Codex both work.

3. Type one sentence: **"Read START-HERE.md and build my first wiki."**

The agent asks you a few questions, builds the folders, and walks you through your first sources. Give it an afternoon. By the end you'll have a cited answer you didn't have to work out from scratch.

Once you've filed your first answer back, tell me what it showed you at info@arcqusgroup.com.

Source & Attribution

This primer explains and builds on a pattern introduced by Andrej Karpathy. He published it as a public idea file titled LLM Wiki, sharing it not as software but as a written concept, on the view that in an era of capable AI agents the idea itself is what is worth passing along, since each reader's own agent can build a version tailored to their work.

His memorable one-line framing of the setup runs: "Obsidian is the IDE; the LLM is the programmer; the wiki is the codebase."

In brief, the original argues that most people use AI on their documents through retrieval, where the model re-reads source fragments on every question and builds nothing that lasts. The pattern instead has the AI compile sources once into a persistent, interlinked set of notes and keep that synthesis current, so understanding compounds rather than resetting to zero. It describes three layers, the immutable raw sources, the maintained wiki, and a schema that tells the agent how to work, alongside three core operations, ingesting a source, querying the wiki, and periodically checking it for contradictions and gaps.

The complete original text is available from Karpathy's own publication at the link below.

https://gist.github.com/karpathy/442a6bf555914893e9891c11519de94f

This document is an Arcqus.AI explainer of that pattern, written so the core idea is graspable without any prior background. The concept originated with Andrej Karpathy, credited here; for the complete original, follow the link to the source.